

Working with relational databases

Why use a database?

- ▶ Database systems (DBS) are a very popular way of storing data. They are particularly efficient in dealing with large data sets.
- ▶ Think of a database (DB) as a dataset managed by a software (known as the database server). In order to access the data, you have to "talk" to the database server. The same way that web servers managed internet resources.
- ▶ As such, DB system facilitates the sharing of the same data between multiple entities. This reduces the number of copies of the data that is kept.

Why use a database?

- ▶ DBS make it easy to control who can access the data.
- ▶ DBS Allow concurrent access to data - synchronize access.
Can handle transactions. Enforce the security and integrity of the data.
- ▶ Makes it easy to maintain the data.
- ▶ DB systems can accommodate very large data sets. They are highly optimized for fast retrieval of the data.

What do we need to know?

Goals:

- ▶ Develop familiarity with relational database systems, and be able to map simple database terminology into statistics terms.
- ▶ Use the SELECT command of the Structured Query Language (SQL) to access data.

Relational databases

- ▶ Currently, the dominant model for database management is called "relational databases" and we will see why shortly.
- ▶ Major Relational database systems includes **Oracle**, **SQL Server**, **PostgreSQL**, **MYSQL**, **MS Access**, **SQLite**.
- ▶ **Oracle** (Oracle), **SQL Server** (MS), and **MS Access** (MS) are commercial systems.
- ▶ **SQLite**, **MYSQL**, and **PostgreSQL** are open-source. **SQLite** is a light-weight system. **PostgreSQL** is comparable to Oracle.
- ▶ There is a language for manipulating relational databases known as **SQL** (Structured Query Language). We will only learn the fraction of SQL that deals with accessing the data (the **SELECT** command).

Relational databases

- ▶ In a relational DB system, the data is represented in two dimensional **tables** (also called relations) which consist of rows (records, observations) and columns (attributes or variables). A database is a collection of such tables.
- ▶ Think of a table as a dataframe that is holding variables (or column or attributes). Each row of the table is an observation (or a case, or tuple).

Relational databases

- ▶ In DB systems, effort is put in representing the data in a way that avoids redundancy as much as possible.
- ▶ This is done by breaking the original data across multiple tables. A process known as **normalization**.
- ▶ This is best understood through an example.

Relational databases

- ▶ Suppose that you are managing the data for a company with a number of customers, each of which will be placing multiple orders, each order has multiple items.
- ▶ If we try to keep one table to record the data, it will look like this.

CNo	Name	Address	ONo	OrderDate	ItemNo	ItemName	ItemQty	Price
1	Smith, F	101 Elm	201	"1-12-99"	001	Item1	100	5
1	Smith, F	101 Elm	201	"1-12-99"	005	Item5	250	10
2	Brown, D	17 Spruce	301	"08-27-01"	010	Item10	20	350
2	Brown, D	17 Spruce	302	"02-21-04"	031	Item31	400	12

Relational databases

- ▶ Each time an order is placed you will need to repeat the customer information, including the Customer's name, address etc...
- ▶ Worst: for each item, you need to repeat the Order information and the customer information.
- ▶ For example if a customer place two orders each with 4 items, you need 8 lines to record these two transactions.
- ▶ What if the customer send you a change of address notice?
- ▶ Difficulties of this one table model: error prone, difficult to maintain.

Relational databases

- ▶ Instead we split this data into three separate tables.
- ▶ One table for customers. If an existing customer places a new order this table does not need to be changed.
- ▶ One table for orders. Even if a new order has multiple items, that table gets only one new record.
- ▶ And one table for items.

Relational databases

CNo	Name	Address
1	Smith, F	101 Elm
2	Brown, D	17 Spruce

OrderNo	CNo	OrderDate	DelivAddr
201	1	"1-12-1999"	Address
301	2	"08-27-2001"	Address
302	2	"02-21-2004"	Address

OrderDetailNo	OrderNo	ItemName	ItemQty	ItemPrice
001	201	Item1	100	5
005	201	Item5	250	10
010	301	Item10	20	350
031	302	Item31	400	12

Relational databases

- ▶ The columns "CNo", "OrderNo" and "OrderDetailNo" are called **Primary keys**. Each table in a relational database needs one (and only one) Primary key. That is a column or a group of columns that uniquely identify each record of the table.
- ▶ Notice the presence of the column "CNo" (the Primary key of the table Customer) in the table "Order". In table "Order", the column "CNo" is called a **Foreign Key**.
- ▶ In relational databases, by creating foreign keys where needed, one create links between tables.

Relational databases

- ▶ Given a database, as a statistician, we need some basic ability to query the database and retrieve the data.
- ▶ By installing the appropriate additional library/package in R we can work with with most databases from within R.
- ▶ To do so, we need a brief introduction to working with packages in R.

R packages

- ▶ Packages in R are pieces of software that can be loaded into a R session, to acquire additional computing tools.
- ▶ They are written by members of the community and freely available online (www.r-project.org).
- ▶ However, we can install packages from within R, without using a browser.

R packages

- ▶ You can see the list of packages attached to your current R session with the function `.packages()`, as follows.

```
z = .packages()
```

```
z
```

```
[1] "space"      "stats"      "graphics"   "grDevices" "ut
```

```
[6] "datasets"  "methods"    "base"
```

```
>
```

- ▶ You can have a description of a package with `packageDescription("stats")`

R packages

- ▶ Before we can use a package, we need to install and load it.
- ▶ We typically install a package using the function `install.packages`.
- ▶ For instance, we will need below to use the package `RSQLite` that will allow us to connect to SQLite databases. We can do this as follows.

```
install.packages("RSQLite", dep=T)
```

- ▶ The argument `dep` specifies whether we would like to install at the same time all other packages on which `RSQLite` is built (in this case, the package `RSQLite` needs another important package called `DBI`).

R packages

- ▶ When a package is installed, we can load it into the current R session by calling the function `library` as follows.

```
library(RSQLite)# or library("RSQLite")
```

- ▶ You can also use the function `require`.
- ▶ When done with a Package, you can remove it from the search directory with the function `detach`, as in

```
detach(package:RSQLite)
```