

Statistics 700 Homework 3 (alpha)

Based on solution by Po-Heng Chen

October 15, 2017

ESS for Importance Sampling

1. Implement an importance sampler and calculate ESS(m) for $m = 50, 100, 200, 500, 1000$.

Suppose $\pi(x)$ is the target distribution where $\pi(x) \propto N(0, 1)$, and $g(x)$ is the importance function (proposal distribution) where $g(x) \propto$ Student's t distribution with 2 degrees of freedom.

Let's first try to build the importance sampler with sample size $m = 50, 100, 200, 500, 1000$.

```
for( m in c(50, 100, 200, 500, 1000)){
  set.seed(700)
  a = 10
  target = function(x) {dnorm(x, mean = 0, sd = 1)}
  n = m
  # importance function: student t with 2 df
  rs = rt(n, df = 2)
  g = dt(rs, df = 2)
  q = target(rs)
  w = q / g
  rstarget = sample(rs, size = n, prob = w, replace = TRUE)
  hist(rstarget, 30, freq = FALSE,
       main = paste0('Recover target density (' ,m, ')'))
  x = seq(-a, a, length.out = 200)
  c = sum(target(x) * (x[2] - x[1]))
  lines(x, target(x)/c, col = 'red', lwd = 2)
  ess = m / (1 + var(w)/(mean(w))^2)
  print(paste0('ESS(' , m, ') = ' , ess))
}
```

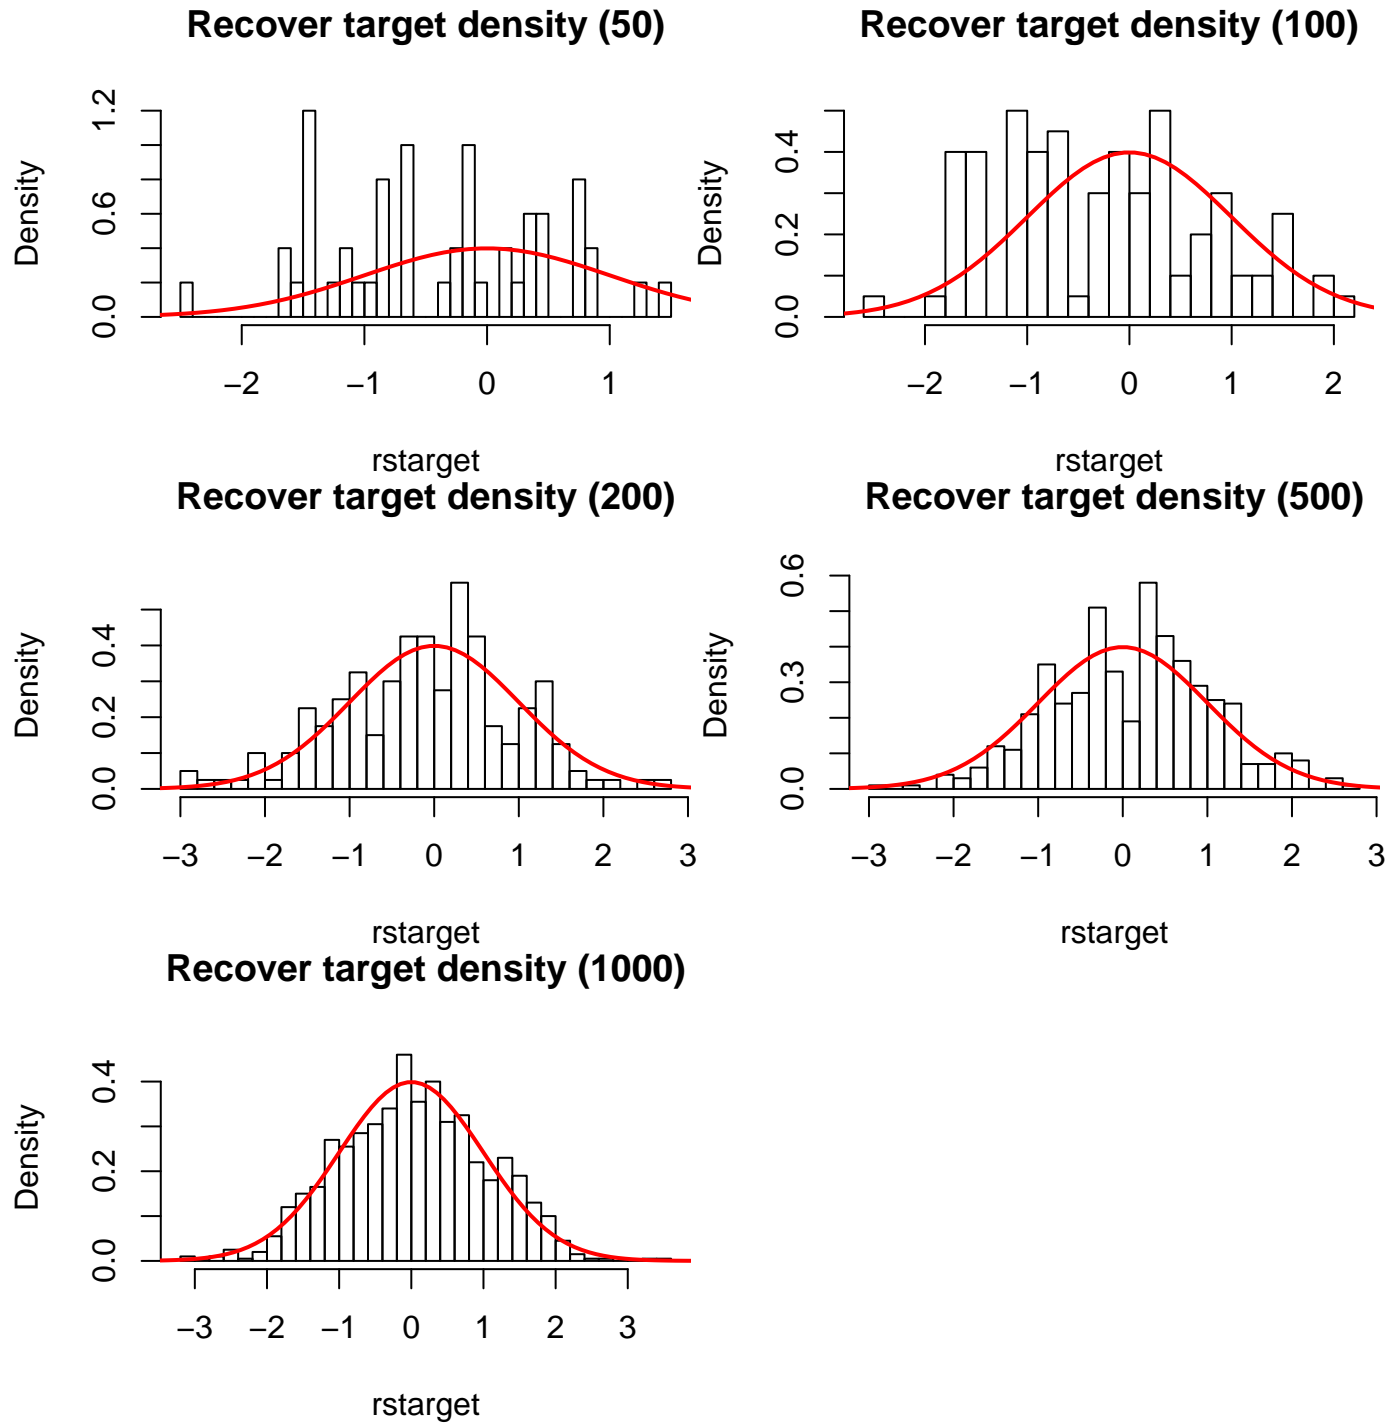
```
## [1] "ESS(50) = 46.8008333894794"
```

```
## [1] "ESS(100) = 89.6052735109369"
```

```
## [1] "ESS(200) = 176.12806971714"
```

```
## [1] "ESS(500) = 432.447278250017"
```

```
## [1] "ESS(1000) = 863.719473712893"
```



2. Implement an importance sampler and calculate $ESS(m)$ for $m = 50, 100, 200, 500, 1000$.

Since the target function and the importance function are switched, we hence switch target/importance functions in the above code.

```
for(m in c(50, 100, 200, 500, 1000)){
  set.seed(700)
  a = 10
  target = function(x) {dt(x, df = 2)}
```

```

n = m
# importance function: standard normal density
rs = rnorm(n, mean = 0, sd = 1)
g = dnorm(rs, mean = 0, sd = 1)
q = target(rs)
w = q / g
rstarget = sample(rs, size = n, prob = w, replace = TRUE)
hist(rstarget, 30, freq = FALSE,
     main = paste0('Recover target density (' ,m, ')'))
x = seq(-a, a, length.out = 200)
c = sum(target(x) * (x[2] - x[1]))
lines(x, target(x)/c, col = 'red', lwd = 2)
ess = m / (1 + var(w)/(mean(w))^2)
print(paste0('ESS(' , m, ') = ' , ess))
}

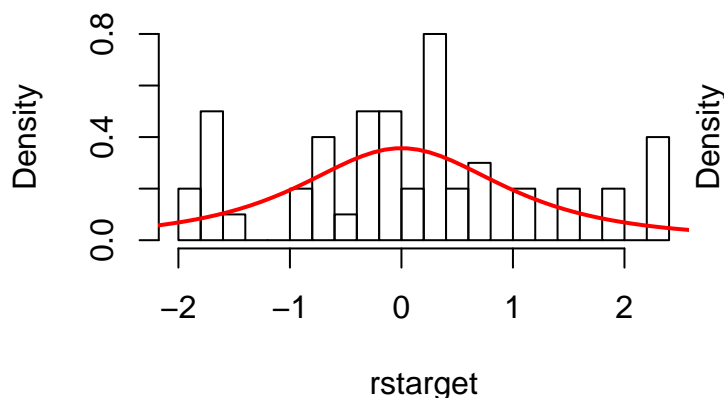
```

```


## [1] "ESS(50) = 48.3337168262753"
## [1] "ESS(100) = 77.2077131716916"
## [1] "ESS(200) = 155.836641958942"
## [1] "ESS(500) = 341.06417912612"
## [1] "ESS(1000) = 733.985499433072"

```

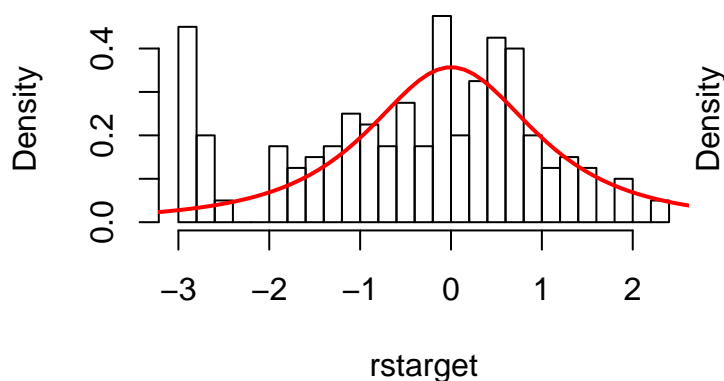
Recover target density (50)



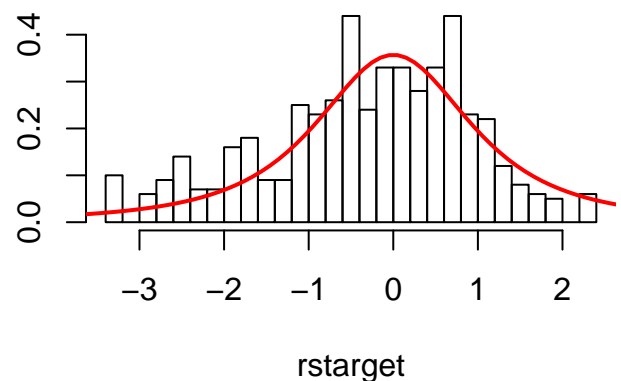
Recover target density (100)




Recover target density (200)



Recover target density (500)




Recover target density (1000)



However, ESS in this case is extremely unstable, since it has infinite variance.

```
ess <- c()
m.seq <- c(50, 100, 200, 500, 1000)
replicant <- function(){
  for(i in 1:length(m.seq)){
    m <- m.seq[i]
    a = 10
    target = function(x) {dt(x, df = 2)}
    n = m
    # importance function: student t with 2 df
    rs = rnorm(n, mean = 0, sd = 1)
    g = dnorm(rs, mean = 0, sd = 1)
    q = target(rs)
    w = q / g
    ess[i] = m / (1 + var(w)/(mean(w))^2)
  }
  ess
}
ess <- replicant()
plot(m.seq, ess, type = 'b', ylim = c(0, 1000), las = 1)
for (i in 1:5) {
  ess <- replicant()
  lines(m.seq, ess, type = 'b')
}
```



3. What do you find by comparing the results above?

Using a heavy tailed distribution as proposal distribution for a (sub-)Gaussian distribution produces weights with infinite variance, and leads to unstable estimates.


4.

(a) Contour Plot


It's important that we truncate it!

```
target <- function(para){
  mu = para[1]
  var = para[2]
  if ((mu < 5) && (mu > -3) && (var > 0.01) && (var < 50)) {
    return(var^(-5/2)*exp(-((mu-1)^2 + 4)/(2*var)))
  } else {
    return(0)
  }
}

xseq = seq(-3, 5, length.out = 100)
yseq = seq(0.01, 5, length.out = 100)
z = matrix(0, length(xseq), length(yseq))
for(i in 1:length(xseq)){
  for(j in 1:length(yseq)){
    z[i, j] = target(c(xseq[i], yseq[j]))
  }
}
contour(xseq, yseq, z)
```




```
persp(xseq, yseq, z, theta=15, phi=0)
```



(b)

Let's start with flat importance function of uniform $[0, 4]$. (This is not sampling on a grid, but it's fine.)

```
# flat importance function
Nsamples = 1000
samplesunif = cbind(runif(Nsamples, min = -3, max = 5), runif(Nsamples, min = 0.01, max = 50))
w = apply(samplesunif, 1, target)
idx = sample(1:length(w), replace = TRUE, prob = w/sum(w))
sampletarget2d = samplesunif[idx, ]
rf <- colorRampPalette(rev(brewer.pal(11,'Spectral'))))
h <- hexbin(samplertarget2d)
plot(h, xlab = 'x', ylab = 'y')
```



```
print(Nsamples / (1 + var(w)/(mean(w))^2))
```

```
## [1] 63.59397
```

We can see that the ESS now is around 64.

Next we use a Gaussian distribution as proposal (Note in general it's NOT OK to use a proposal distribution with lighter tails! We allow such a choice only because the target distribution is truncated.)


```
# importance function
g <- function(para){
  mu1 = c(1.5,1.5)
  Sigma1 = sqrt(c(1,2))
  x1 = sum(dnorm(para, mu1, Sigma1, log = TRUE))
  return (exp(x1))
}
```

```

}
xseq = seq(-3, 5, length.out = 100)
yseq = seq(0.01, 5, length.out = 100)

z = matrix(0, length(xseq), length(yseq))
for(i in 1:length(xseq)){
  for(j in 1:length(yseq)){
    z[i, j] = g(c(xseq[i], yseq[j]))
  }
}
contour(xseq, yseq, z)


```



```

persp(xseq, yseq, z, theta=10, phi=0)

```

```
# new importance function
for( m in c(50, 100, 200, 500, 1000)){
  samplesG = cbind(rnorm(Nsamples, 1, sqrt(1)), rnorm(Nsamples, 1, sqrt(1)))

  w = apply(samplesG, 1, target) / apply(samplesG, 1, g)
  w = na.omit(w)
  idx = sample(1:length(w), replace = TRUE, prob = w/sum(w))
  sampletarget = samplesG[idx, ]
  ess = m / (1 + var(w)/(mean(w))^2)
  print(paste0('ESS(', m, ') = ', ess))
}
```

```
## [1] "ESS(50) = 18.7911885614051"
## [1] "ESS(100) = 24.8506648473961"
## [1] "ESS(200) = 45.1006844268438"
## [1] "ESS(500) = 64.3423438854661"
## [1] "ESS(1000) = 523.609275588432"
```

A much better ESS.

Rejection Control Algorithm

Suppose $\pi(x) \propto$ Gaussian density with mean 0 and standard deviation 0.3, $g(x) \propto$ student t distribution with 2 degrees of freedom.

```
rejectioncontrol <- function(Nsamples, samplefromtrial, logtarget, logtrial,
                             c, dimension = 1){
  samplesX <- samplefromtrial(Nsamples)
  if(dimension == 1){
```

```

weights <- sapply(samplesX, function(x) logtarget(x) - logtrial(x))
}
if(dimension > 1){
weights <- as.numeric(apply(samplesX, 1, function(x) logtarget(x) - logtrial(x)))
}
logc <- -Inf
if(c > 0){
logc <- log(c)
}
acceptrates <- rep(1, Nsamples)
indicator <- which(weights < logc)
if(length(indicator) > 0){
acceptrates[indicator] <- exp(weights[indicator] - logc)
}
acceptindicator <- which(runif(Nsamples) < acceptrates)
if(dimension == 1){
acceptedsamples <- samplesX[acceptindicator]
}
if(dimension > 1){
acceptedsamples <- samplesX[acceptindicator, ]
}
weightsnew <- weights[acceptindicator] - log(acceptrates[acceptindicator])
weightsnew <- exp(weightsnew - max(weightsnew))
weightsnew <- weightsnew / sum(weightsnew)
return(list(samples = acceptedsamples, weights = weightsnew))
}

```

1. Implement the importance sampling with and without rejection control.


Implementation without rejection control.

```

m = 2000
set.seed(700)
a = 10
target = function(x) {dnorm(x, mean = 0, sd = 0.3)}
n = m
# importance function: student t with 2 df
rs = rt(n, df = 2)
g = dt(rs, df = 2)
q = target(rs)
w = q / g
rtarget = sample(rs, size = n, prob = w, replace = TRUE)
hist(rtarget, 30, freq = FALSE, main = 'Recover target density')
x = seq(-a, a, length.out = 200)
c = sum(target(x) * (x[2] - x[1]))
lines(x, target(x)/c, col = 'red', lwd = 2)

```

Recover target density



```
ess = m / (1 + var(w)/(mean(w))^2)
print(paste0('ESS(', m, ') = ', ess))
```

```
## [1] "ESS(2000) = 722.131849768128"
```

Implementation with rejection control.

```
samplefromtrial <- function(n){
  rt(n, df = 2)
}
logtrial <- function(x){
  dt(x, df = 2, log = TRUE)
}
logtarget <- function(x){
  dnorm(x, mean = 0, sd = 0.3, log = TRUE)
}
Nsamples <- 2000
c = 5
RCresult <- rejectioncontrol (Nsamples, samplefromtrial, logtarget,
                             logtrial, c, dimension = 1)

w = RCresult$weights
m = length(w)
ess = m / (1 + var(w)/(mean(w))^2)
print(paste0('ESS(', m, ') = ', ess))
```

```
## [1] "ESS(411) = 411"
```

Rejection control provides with LESS effective samples! However this saves computation, as we have to evaluate h on much less samples.

2. Do you see improvement by using rejection control with different C ?

```
for(c in c(0.01, 0.1, 1, 2, 3, 5, 10)){
  RCresult <- rejectioncontrol (Nsamples, samplefromtrial, logtarget,
                              logtrial, c, dimension = 1)

  w = RCresult$weights
  m = length(w)
  ess = m / (1 + var(w)/(mean(w))^2)
  print(paste0('c = ', c, '; ESS(', m, ') = ', ess))
}
```

```
## [1] "c = 0.01; ESS(1265) = 709.58223838829"
## [1] "c = 0.1; ESS(1143) = 745.71178400594"
## [1] "c = 1; ESS(905) = 754.685422757226"
## [1] "c = 2; ESS(736) = 692.846226062537"
## [1] "c = 3; ESS(646) = 640.439335366042"
## [1] "c = 5; ESS(385) = 385"
## [1] "c = 10; ESS(194) = 194"
```

No.

3. What value of c do you choose to use finally and why?


In this implementation, the smaller the better. Since we always discard samples. When $c \rightarrow 0$, we are back in the Importance Sampling case.

4. Repeat the above procedures when $\pi(x) \propto$ Gaussian density with mean 0 and standard deviation 3.

Without rejection control.

```
set.seed(700)
a = 10
target = function(x) {dnorm(x, mean = 0, sd = 3)}
n = m
# importance function: student t with 2 df
rs = rt(n, df = 2)
g = dt(rs, df = 2)
q = target(rs)
w = q / g
rtarget = sample(rs, size = n, prob = w, replace = TRUE)
hist(rtarget, 30, freq = FALSE, main = 'Recover target density')
x = seq(-a, a, length.out = 200)
c = sum(target(x) * (x[2] - x[1]))
lines(x, target(x)/c, col = 'red', lwd = 2)
```

Recover target density



```
ess = m / (1 + var(w)/(mean(w))^2)
print(paste0('ESS(', m, ') = ', ess))
```

```
## [1] "ESS(194) = 95.4471120224163"
```

With rejection control and different C .

```
samplefromtrial <- function(n){
  rt(n, df = 2)
}
logtrial <- function(x){
  dt(x, df = 2, log = TRUE)
}
logtarget <- function(x){
  dnorm(x, mean = 0, sd = 3, log = TRUE)
}
Nsamples <- 2000
for(c in c(0.01, 0.1, 1, 2, 3, 5, 10)){
  RCresult <- rejectioncontrol (Nsamples, samplefromtrial, logtarget,
                                logtrial, c, dimension = 1)

  w = RCresult$weights
  m = length(w)
  ess = m / (1 + var(w)/(mean(w))^2)
  print(paste0('c = ', c, '; ESS(', m, ') = ', ess))
}

## [1] "c = 0.01; ESS(1996) = 982.844581618895"
## [1] "c = 0.1; ESS(1985) = 997.496787312193"
## [1] "c = 1; ESS(1293) = 920.278698503296"
## [1] "c = 2; ESS(827) = 751.096901982948"
## [1] "c = 3; ESS(600) = 587.822124571759"
## [1] "c = 5; ESS(412) = 412"
## [1] "c = 10; ESS(199) = 199"
```